



Temporal Difference Learning

KENNETH TRAN

Principal Research Engineer, MSR AI





Temporal Difference Learning


Policy Evaluation

Outline

- Intro to model-free learning
- Monte Carlo Learning
- Temporal Difference Learning
- TD(λ)

Revisit notations

- Episode



- Return

Discount rate, e.g. 0.9

$$\begin{aligned} G_t &\stackrel{\text{def}}{=} R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \dots) \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

- state-value function

True value of state s under policy π

$$\begin{aligned} v_{\pi}(s) &\stackrel{\text{def}}{=} \mathbf{E}_{\pi}[G_t | S_t = s] \\ &= \mathbf{E}_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbf{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \end{aligned}$$

Recap: Dynamic Programming

- Bellman equation

$$Q^*(s, a) = \sum_{s'} P(s'|s, a) \left(R(s, a, s') + \gamma \max_{a'} Q^*(s', a') \right)$$

- DP: solve a known MDP

Policy Evaluation

Monte Carlo Learning



Overview

- MC methods learn from episodes of experience
- MC is model-free: no knowledge of MDP equations
- MC learns from complete episodes
- MC uses the simplest possible idea: use empirical mean to approximate the expected value
- Caveat: can only apply MC to episodic MDPs
(i.e. all episodes must terminate)

MC for Policy Evaluation

- Goal: learn v_π from episodes of experience under π

$$S_1, A_1, R_2, \dots, S_k$$

- Recall: return is the total discounted reward

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- Recall: value function is the expected return

$$v_\pi(s) = \mathbf{E}[G_t | S_t = s]$$

- MC policy evaluation uses *empirical mean* return instead of *expected* return

MC for Policy Evaluation

To evaluate $v_\pi(s)$

- Sample episodes of experience under π
- Every time t that state s is visited in an episode
 - Increment counter $N(s) \leftarrow N(s) + 1$
 - Increment total return $S(s) \leftarrow S(s) + G_t$
 - Value is estimated by mean return $V(s) = S(s)/N(s)$
- By law of large numbers: $V(s) \rightarrow v_\pi(s)$ as $N(s) \rightarrow \infty$

Incremental MC Updates

- Update $V(s)$ incrementally after episode

$$S_1, A_1, R_2, \dots, S_T$$

- For each state s , with return G_t
 - $N(S_t) \leftarrow N(S_t) + 1$
 - $V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$

Policy Evaluation

Temporal Difference Learning



Temporal Difference Learning

- Like MC, TD is model-free: no knowledge of MDP transition/rewards
- Unlike MC, TD learns from incomplete episodes, by *bootstrapping*
- TD updates a guess towards a guess


MC and TD

- Goal: policy evaluation

For a given policy v_π , compute the state-value function v_π

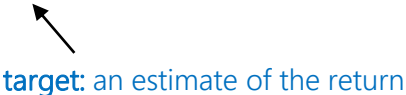
- Recall: simple every-visit Monte Carlo method

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$


step-size parameter target: the actual return after time t

- Simplest temporal-difference learning method: TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$


target: an estimate of the return


$\delta_t \stackrel{\text{def}}{=} R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called the *TD error*

Example: Driving Home


State	Elapsed Time (minutes)	Predicted Time to Go	Predicted Total Time
leaving office	0	30	30
reach car, raining	5	35	40
exit highway	20	15	35
behind truck	30	10	40
home street	40	3	43
arrive home	43	0	43

Driving home example: MC vs TD

Changes recommended by
Monte Carlo methods ($\alpha=1$)



Changes recommended
by TD methods ($\alpha=1$)



Pros and Cons of MC vs. TD (1)

- TD can learn *before* knowing the final outcome
 - TD can learn online after every step
 - MC must wait until end of episode before return is known
- TD can learn *without* the final outcome
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments



Bias/Variance Trade-Off

- Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is *unbiased* estimate of $v_\pi(S_t)$
- TD target $R_{t+1} + \gamma V(S_{t+1})$ is *biased* estimate of $v_\pi(S_t)$


Pros and Cons of MC vs. TD (2)

- MC has high variance, zero bias
 - Not very sensitive to initial value
 - Very simple to understand and use
- TD has low variance, some bias
 - Usually more efficient than MC
 - More sensitive to initial value

Random Walk Example



Random Walk: MC vs. TD



Batch MC and TD

- MC and TD converge: $V(s) \rightarrow v_\pi(s)$ as experience $\rightarrow \infty$
- But what about batch solution for finite experience?

$$\begin{aligned} & s_1^1, a_1^1, r_2^1, \dots, s_{T_1}^1 \\ & \dots \\ & s_1^K, a_1^K, r_2^K, \dots, s_{T_1}^K \end{aligned}$$

- e.g. Repeatedly sample episode $k \in [1, K]$
- Apply MC or TD(0) to episode k
- Will they converge to the same value function?

AB Example

Two states A , B ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$


$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$



What is $V(A), V(B)$?

Certainty Equivalence

- MC converges to solution with minimum mean-squared error
 - Best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} \left(G_t^k - V(s_t^k) \right)^2$$

- In the AB example, $V(A) = 0$
- TD(0) converges to solution of max likelihood Markov model
 - Solution to the MDP $\langle S, A, P, R, \gamma \rangle$ that best fits the data
 - In the AB example, $V(A) = 0.75$

Pros and Cons of MC vs. TD (3)

- TD exploits Markov property → more efficient
- MC does not exploit Markov property

Unified View of

- Exhaustive Search
- Dynamic Programming
- Monte Carlo
- and Temporal Difference Learning

Monte-Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$




Image Credit: David Silver, Model-Free Prediction, UCL Course on RL

Temporal-Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$




Image Credit: David Silver, Model-Free Prediction, UCL Course on RL

Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1})]$$










Image Credit: David Silver, Model-Free Prediction, UCL Course on RL

Bootstrapping and Sampling

- **Bootstrapping:** update involves an *estimate*
 - MC does not bootstrap: it learns from complete episodes 
MC must wait until end of episode before return is known
 - DP bootstrap 
 - TD bootstrap 
- **Sampling:** update does not require computing exact *expectation*
 - MC samples 
 - DP does not sample 
 - TD samples 

Unified View of Reinforcement Learning




Image Credit: David Silver, Model-Free Prediction, UCL Course on RL

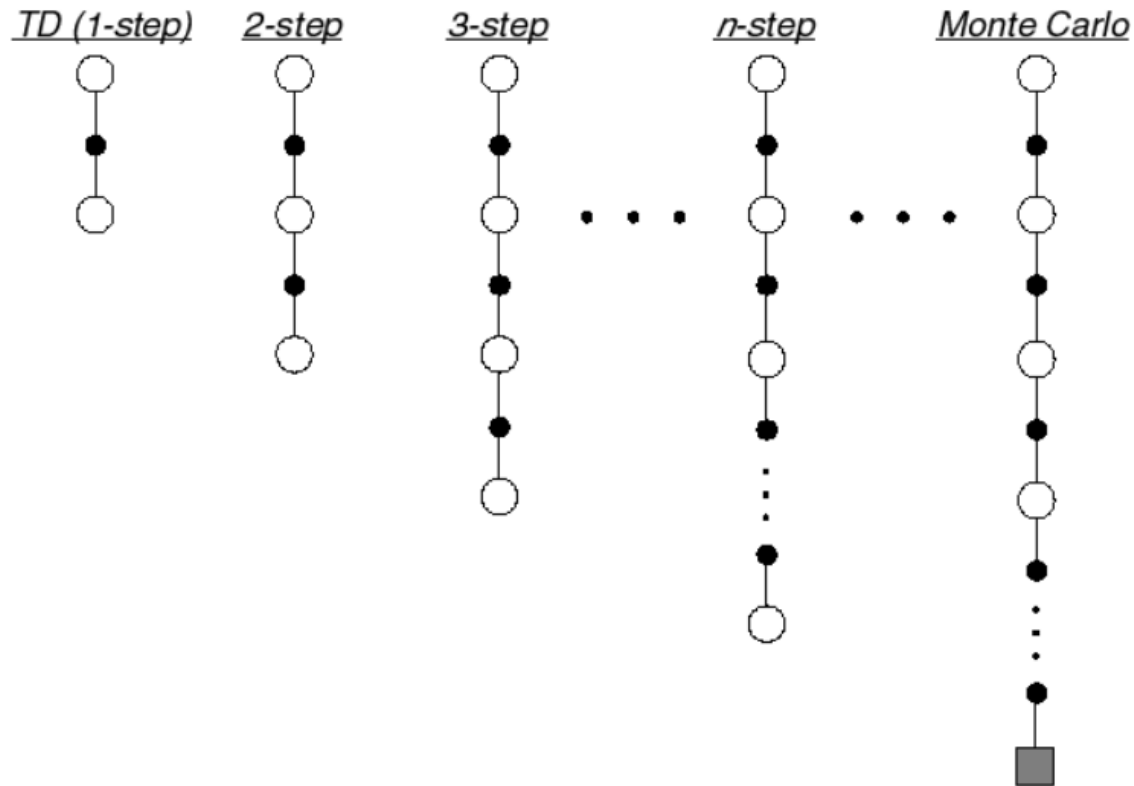
Policy Evaluation

$TD(\lambda)$



n-Step Prediction

Let TD target look n steps into the future



n-Step Return

- Consider the following n -step returns for $n = 1, 2, \infty$:

$$n = 1 \quad (TD) \quad G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$$

$$n = 2 \quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma V(S_{t+2})$$

...

$$n = \infty \quad (MC) \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{T-1} R_T$$

- Define the n -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- n -step temporal-difference learning

$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{(n)} - V(S_t) \right)$$

Large Random Walk Example

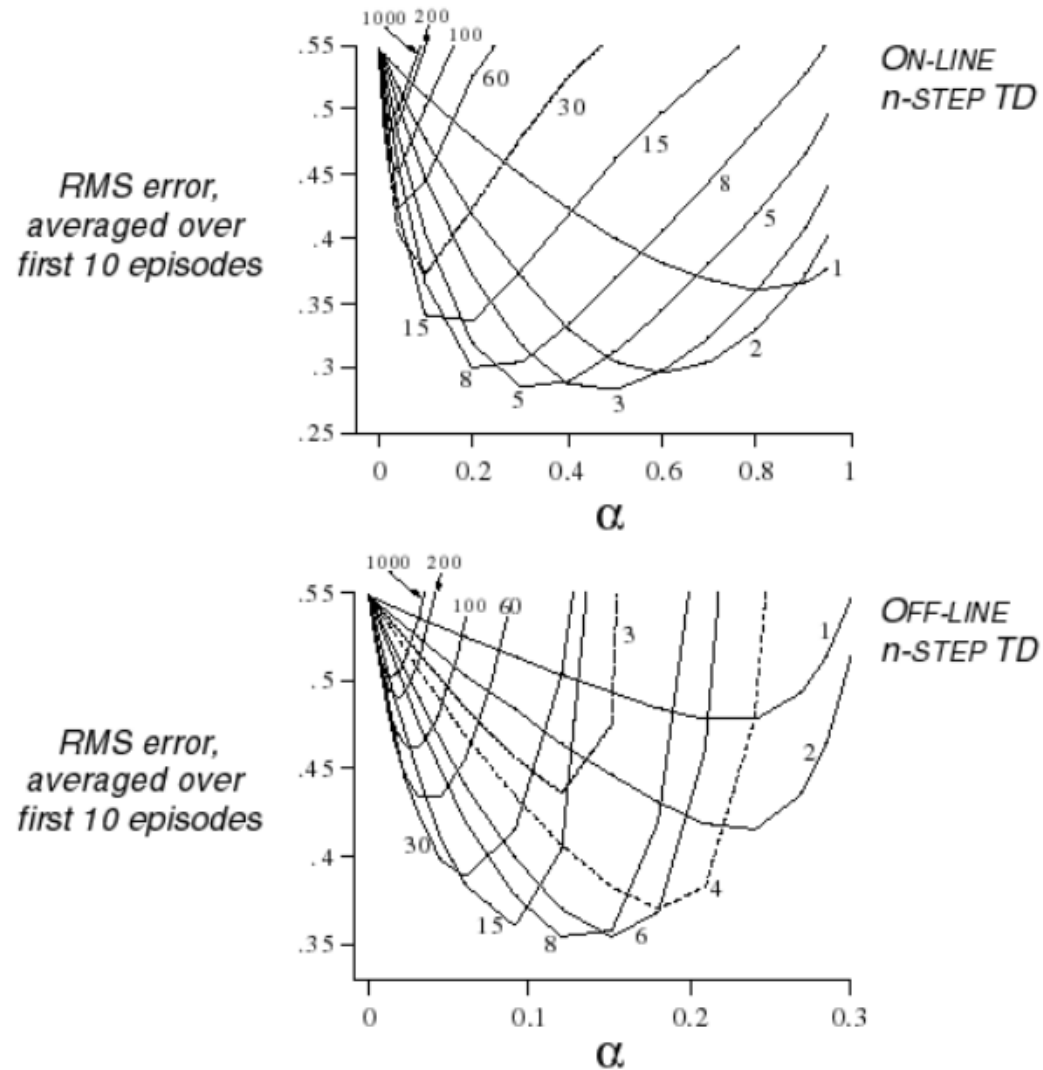



Image Credit: Sutton and Barto, Reinforcement Learning, An Introduction 2017

Averaging n -Step Returns

- We can average n -step returns over different n
- e.g. average the 2-step and 4-step returns

$$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)}$$

- Combines information from two different time-steps
- Can we efficiently combine information from all time-steps?




λ -return

- The λ -return G_t^λ combines all n -step returns G_t^n
- Using weight $(1 - \lambda)\lambda^{n-1}$


$$G_t^\lambda = (1 - \lambda)$$

- Forward-view TD(λ)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$




TD(λ) Weighting Function



$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

Forward-view TD(λ)




- Forward-view looks into the future to compute G_t^λ
- Like MC, can only be computed from complete episodes

Backward View TD(λ)

- Forward view provides theory
- Backward view provides mechanism
- Update online, every step, from incomplete sequences


Eligibility Traces



- Credit assignment problem: did bell or light cause shock?
- **Frequency heuristic**: assign credit to most frequent states
- **Recency heuristic**: assign credit to most recent states
- *Eligibility traces* combine both heuristics

$$E_0(s) = 0$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$



TD(λ) Algorithm

At each time step t in the rollout

- Update eligibility trace for every state s


$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

current state

- Compute TD-error $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$

- Update value function $V(s)$ for every state s

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



TD(λ) and TD(0)

- When $\lambda = 0$, only current state is updated

$$E_t(s) = \mathbf{1}(S_t = s)$$
$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$

- This is exactly equivalent to TD(0) update

$$V(S_t) \leftarrow V(S_t) + \alpha \delta_t$$

TD(λ) and MC

- TD(1) is roughly equivalent to MC
- Error is accumulated online, step-by-step
- If value function is only updated offline at end of episode
- Then total update is exactly the same as MC

Summary of Forward & Backward TD(λ)

- Batch Update
 - Updates are accumulated within episode, but applied in batch at the end of episode
 - \rightarrow Backward TD(λ) is equivalent to Forward TD(λ)
- Online Update
 - TD(λ) updates are applied at each step within episode
 - Forward and backward-view TD(λ) are slightly different
 - Backward TD(λ) is typically more efficient
 - Analogy: SGD vs. batch GD



Temporal Difference Learning

Policy Optimization

Outline

- Introduction
- On-policy Monte Carlo control
- On-policy Temporal Difference control
- Off-policy Learning
- Summary


Introduction

- Last lesson: model-free **policy evaluation**
 - Estimate the value function of an *unknown* MDP
- This lesson: model-free **policy optimization**
 - Optimize the value function of an *unknown* MDP

On and Off-Policy Learning



- On-policy learning
 - “Learn on the job”
 - Learn about policy π from experience sampled from π
- Off-policy learning
 - “Look over someone’s shoulder”
 - Learn about policy π from experience sampled from μ

Generalised Policy Iteration (Refresher)



Policy evaluation Estimate v_π
e.g. Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$
e.g. Greedy policy improvement




Policy Optimization

On-Policy MC Control



Generalised Policy Iteration With Monte-Carlo Evaluation



Policy evaluation Monte-Carlo policy evaluation, $V = v_\pi$?

Policy improvement Greedy policy improvement?

Model-Free Policy Iteration Using Action-Value Function


- Greedy policy improvement over $V(s)$ requires model of MDP

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s')$$

- Greedy policy improvement over $Q(s, a)$ is model-free

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$$

Generalised Policy Iteration with Action-Value Function



Policy evaluation Monte-Carlo policy evaluation, $Q = q_\pi$

Policy improvement Greedy policy improvement?

Example of Greedy Action Selection



"Behind one door is tenure - behind the other is flipping burgers at McDonald's."

- There are two doors in front of you.
- You open the left door and get reward 0 $V(\text{left}) = 0$
- You open the right door and get reward +1: $V(\text{right}) = +1$
- You open the right door and get reward +3: $V(\text{right}) = +3$
- You open the right door and get reward +2: $V(\text{right}) = +2$
- ...
- Are you sure you've chosen the best door?


ϵ -Greedy Exploration

- Simplest idea for ensuring continual exploration
- All m actions are tried with non-zero probability
- With probability $1 - \epsilon$ choose the greedy action
- With probability ϵ choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon & \text{if } a^* = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

- **Theorem:** For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to Q_π is an improvement, i.e.
 $v_{\pi'}(s) \geq v_\pi(s)$

Monte Carlo Control



Every episode:

Policy evaluation Monte-Carlo policy evaluation, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

GLIE

Definition: Greedy in the Limit with Infinite Exploration (GLIE)

- All state-action pairs are explored infinitely many times

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- The policy converges on a greedy policy

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = \mathbf{1}(a = \operatorname{argmax}_{a' \in A} Q_k(s, a'))$$

- Example: ϵ -greedy is GLIE if ϵ reduces to zero at $\epsilon_k = \frac{1}{k}$

GLIE Monte-Carlo Control

- Sample k^{th} episode using $\pi: \{S_1, A_1, R_2, \dots, S_T\} \sim \pi$
- For each state S_t and action A_t in the episode,

$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t))$$

- Improve policy based on new action-value function

$$\epsilon \leftarrow 1/k$$

$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

- **Theorem:** GLIE Monte-Carlo control converges to the optimal action-value function

$$Q(s, a) \rightarrow q_*(s, a)$$

Policy Optimization


On-Policy TD Control



MC vs. TD Control


- Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
 - Lower variance
 - Online
 - Incomplete sequences
- Natural idea: use TD instead of MC in our control loop
 - Apply TD to $Q(S, A)$
 - Use ϵ -greedy policy improvement
 - Update every time-step

Sarsa for updating action-value function



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

On-Policy Control With Sarsa



Every **time-step**:

Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Sarsa Algorithm for On-Policy Control

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Convergence of Sarsa

Theorem: Sarsa converges to the optimal action-value function

$$Q(s, a) \rightarrow q_*(s, a)$$



under the following conditions:

- GLIE sequence of policies $\pi_t(a|s)$
- Robbins-Monro sequence of step sizes α_t


$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Windy Gridworld Example



standard
moves



king's
moves

- Reward = -1 per time step until reaching goal
- Undiscounted

Sarsa on the Windy Gridworld




Image Credit: Sutton and Barto, Reinforcement Learning, An Introduction 2017

Sarsa(λ)

n -Step Sarsa

- Consider the following n -step returns for $n = 1, 2, \dots, \infty$:

$$n = 1 \quad (\text{Sarsa}) \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$$

$$n = 2 \quad q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}, A_{t+2})$$

...

...

$$n = \infty \quad (MC) \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$


- Define the n -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n}, A_{t+n})$$

- n -step Sarsa updates $Q(s, a)$ towards the n -step Q-return

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left(q_t^{(n)} - Q(S_t, A_t) \right)$$

Forward View Sarsa(λ)



- The q^λ return combines all n -step Q-returns $q_t^{(n)}$
- Using weight $(1 - \lambda)\lambda^{n-1}$

$$q_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

- Forward-view Sarsa(λ)

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (q_t^\lambda - Q(S_t, A_t))$$

Backward View Sarsa(λ)

- Just like TD(λ), we use **eligibility traces** in an online algorithm
- But Sarsa(λ) has one eligibility trace for each state-action pair

$$E_0(s, a) = 0$$

$$E_t(s, a) = \gamma\lambda E_{t-1}(s, a) + \mathbf{1}(S_t = s, A_t = a)$$

- $Q(s, a)$ is updated for every state s and action a
- In proportion to TD-error δ_t and eligibility trace $E_t(s, a)$

$$\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta_t E_t(s, a)$$

Sarsa(λ) Algorithm

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Repeat (for each episode):

$E(s, a) = 0$, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$

Initialize S, A

Repeat (for each step of episode):

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$

$E(S, A) \leftarrow E(S, A) + 1$

For all $s \in \mathcal{S}, a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$


$E(s, a) \leftarrow \gamma \lambda E(s, a)$

$S \leftarrow S'; A \leftarrow A'$


until S is terminal

Sarsa(λ) Gridworld Example


Path taken



Action values increased
by one-step Sarsa



Action values increased
by Sarsa(λ) with $\lambda=0.9$



Policy Optimization

Off-Policy Learning



Off-Policy Learning

- Evaluate target policy $\pi(a|s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$

- While following behaviour policy $\mu(a|s)$

$$\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$$

- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
 - Learn about *optimal* policy while following *exploratory* policy
 - Learn about *multiple* policies while following *one* policy

Off-Policy Control with Q-Learning


- The **target policy** π is **greedy** w.r.t. $Q(s, a)$

$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} Q(S_{t+1}, a')$$

- The **behaviour policy** μ is e.g. **ϵ -greedy** w.r.t. $Q(s, a)$
- The Q-learning target becomes:

$$R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a')$$

Q-Learning Control Algorithm



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

Theorem: *Q-learning control converges to the optimal action-value function*

$$Q(s, a) \rightarrow q_*(s, a)$$

Q-Learning Algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$;

until S is terminal

References

Richard Sutton. Reinforcement Learning: An Introduction (2nd edition).

David Silver's course on Reinforcement Learning:
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>