



# Reinforcement Learning with Function Approximation

**KATJA HOFMANN**

Researcher, MSR Cambridge



# Contents

## **1 Overview**

Representation and Generalization in RL

## **2 Linear Function Approximation**

Learning generalizable value functions

## **3 Lab 1: RL with Linear Function Approximation**

Navigating grid worlds with linear function approximation

## **4 RL with Deep Neural Networks**

Focus on training stability

## **5 Lab 2: RL with Deep Neural Networks**

Extension to non-linear function approximation and continuous tasks

## **6 Extensions and further reading**

Double Q, prioritized replay, et al.


# Overview

# To solve reinforcement learning,

We must overcome 4 fundamental challenges:

- Representation
- Generalization
- Temporal Credit Assignment
- Exploration

Reminder: challenges  
discussed in module 2




# To solve reinforcement learning,

We must overcome 4 fundamental challenges:


- Representation
- Generalization
- Temporal Credit Assignment
- Exploration

# Challenge: Representation


Grid Maze




Go



Atari (Pong)



Cart Pole



# Aim: generalizable representations

- From state (state-action) values to state (state-action) value functions  $V(\phi(s))$  and  $Q(\phi(s, a))$ , with  $k$ -dimensional observations

$$\phi(s) = \begin{pmatrix} \phi_{11} & \cdots & \phi_{1k} \\ \vdots & \ddots & \vdots \\ \phi_{n1} & \cdots & \phi_{nk} \end{pmatrix}$$


- Why can this work?  
Consider high dimensional observations vs lower dimensional underlying state
- Ideal representation: sample efficient training and good generalization to unseen parts of the observation space

# Starting point: define observation space

- Important aspect of problem formulation
- Opportunity to encode expert knowledge
- Trade-off: sample efficient learning vs generality




# Example: observation space for Pong




Example:  $k = 3 \times 9$

indicate if center of each paddle  
and ball is in a given grid cell



Example:  $k = 3 \times 2$

indicate x, y coordinate of ball  
and paddles



Example:  $k = (84 \times 84)256$

84x84 8-bit gray-scale pixel values  
of scaled image

# Linear Function Approximation

# Linear function approximators

- Recall – recursive formulation of the value function:

$$\begin{aligned} v^\pi(s_t) &= \sum_{s_{t+1}} R(s_t) + P(s_t, s_{t+1}; \pi) \gamma v^\pi(s_{t+1}) \\ &= \mathbb{E}_\pi [R(s_t) + \gamma v^\pi(s_{t+1})] \end{aligned}$$

- And assume linear model:

$$v^\pi(s_t) = \theta^T \phi(s_t)$$

where  $\theta$  denotes the model parameters

- Formulate temporal difference error:

$$\delta_{t+1} = R_t + \gamma \theta^T \phi(s_{t+1}) - \theta^T \phi(s_t)$$

# Iterative algorithm

- Convenient update given samples  $s_t, r_t, s_{t+1}$

$$\delta \leftarrow r_t + \gamma \theta^T \phi(s_{t+1}) - \theta^T \phi(s_t)$$

$$\theta \leftarrow \theta + \alpha \delta \phi(s_t)$$

- Approximately minimizes the squared loss (under certain conditions):

$$J(\theta) = \|r_t + \gamma \theta^T \phi(s_{t+1}) - \theta^T \phi(s_t)\|^2$$

# Extension to linear Q-learning

- Approximate state-action instead of actions values and update given samples  $s_t, a_t, r_t, s_{t+1}$  and feature extractor  $\phi(s, a)$

$$\delta \leftarrow r_t + \gamma \max_{a \in A} \theta^T \phi(s_{t+1}, a) - \theta^T \phi(s_t, a_t)$$

- Corresponding to the squared loss:

$$J(\theta) = \left\| r_t + \gamma \max_{a \in A} \theta^T \phi(s_{t+1}, a) - \theta^T \phi(s_t, a_t) \right\|^2$$

- For a complete Q-learning algorithm with linear function approximation, follow by policy improvement step to derive policy  $\pi(s)$

$$\pi(s) \leftarrow \max_{a \in A} \theta^T \phi(s, a)$$

# Summary

- RL with function approximation, aim: representation that can be learned efficiently and generalizes well
- Trade-off between manually constructed and learned components
- RL with Linear function approximation – approximate values using linear weighted feature combinations

# Further Reading

- **Reinforcement Learning – An Introduction** (Sutton and Barto, 2017):
  - <http://www.incompleteideas.net/sutton/book/the-book-2nd.html>
  - Part II: approximate solution methods + references
- **Algorithms for Reinforcement Learning** (Szepesvári, 2010). *Synthesis Lectures on Artificial Intelligence and Machine Learning*:
  - More details on least squares methods (section 3) + overview of convergence results + references

Lab 1



# RL with Deep Neural Networks

# Beyond linear function approximation

Recall Q-learning with linear function approximation minimizes squared loss between current Q and targets. More generally:

$$J(\theta_i) = \left\| r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) - Q(\phi(s_t, a); \theta_i) \right\|^2$$

Update parameters  $\theta$  using stochastic gradient descent (SGD):

$$\nabla_{\theta_i} J(\theta_i) = \mathbb{E}_{s_t, a_t, r_t, s_{t+1} \sim D} \left[ r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) - Q(\phi(s_t, a_t); \theta) \right] \nabla_{\theta_i} Q(\phi(s_t, a_t); \theta)$$

# Deep Q-Learning: algorithm outline

- Policy evaluation step:

$$\delta_t \leftarrow r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) - Q(\phi(s_t, a); \theta_i)$$

$$\theta_{t+1} = \theta_t + \alpha(\delta_t - Q(\phi(s, a); \theta)) \nabla_{\theta_t} Q(\phi(s, a); \theta_t)$$

- Policy improvement step:

$$\pi_t(s) \leftarrow \max_{a \in A} Q(\phi(s, a); \theta_t)$$

- In practice: use standard deep learning frameworks to compute gradients and apply parameter updates

# Beyond linear function approximation

Stochastic gradient methods – key to using powerful function approximators, in particular deep neural nets

Benefit: use tools from supervised learning – deep learning frameworks, optimizers

Challenge: stability / robustness

# Target Network

In deep Q-learning, regression targets are non-stationary:

$$J(\theta_i) = \left\| r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-) - Q(\phi(s_t, a); \theta_i) \right\|^2$$

Achieve temporary stability, by using a separate target network  $\theta_i^-$  to compute update targets

# Replay Memory

Updates use minibatches sampled from a data buffer  $D$

$$\nabla_{\theta_i} J(\theta_i) = \mathbb{E}_{\mathbf{s}_t, \mathbf{a}_t, r_t, \mathbf{s}_{t+1} \sim D} \left[ r_t + \gamma \max_{a \in A} Q(\phi(\mathbf{s}_{t+1}, a); \theta_i^-) - Q(\phi(\mathbf{s}_t, \mathbf{a}_t); \theta) \right] \nabla_{\theta_i} Q(\phi(\mathbf{s}_t, \mathbf{a}_t); \theta)$$

Implemented as large capacity “replay buffer” of recent interactions

# Summary

- Deep Q-learning formulation enables use of powerful function approximators, e.g., Deep Neural Networks
- Leverage tools from supervised learning for effective learning
- Challenge: no convergence guarantee, need to stabilize learning

# Further Reading

- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533.
  - Introduce DQN, breakthrough result on learning to play Atari from pixels
- Riedmiller, M. (2005). Neural Fitted Q Iteration – First Experiences with a Data Efficient Neural Reinforcement Learning Method. ECML 2005, 317-328.
  - Propose use of neural nets for Q-learning, Rprop optimizer – promising results



# Lab 2

# Extensions and Practical Considerations

# Extensions of Deep Q-Learning

Focus on increasing robustness, reducing bias and/or variance

Here: Double DQN, prioritized replay memory

# Double DQN

Start from target computation for Deep Q-Learning:

$$Y^{DQN} = r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-)$$

Problem: biased estimator –  $\max_{a \in A}$  and  $Q(\phi(s_{t+1}, a); \theta_i^-)$  rely on the same estimator, parameterized by  $\theta_i^-$ . Proposed algorithm – compute targets using:

$$Y^{DDQN} = r_t + \gamma Q(\phi(s_{t+1}, \operatorname{argmax}_{a \in A} Q(\phi(s_{t+1}, a); \theta_i)); \theta_i^-)$$

# Prioritized Replay Memory

Idea – instead of uniform sampling from replay buffer  $D$ , prioritize transitions with high loss.

Proposed approach – sample with

$$p(s_t, a_t, r_t, s_{t+1}) \propto r_t + \gamma \max_{a \in A} Q(\phi(s_{t+1}, a); \theta_i^-)$$

# Practical considerations

For SGD / network architecture – see supervised deep learning approaches

Additional consideration for Deep Q-learning: exploration, replay memory size

Data efficiency – interactions may be slow / costly; reference work on multi-task / meta-learning

# Further Reading

- van Hasselt, H., Guez, A., & Silver, D. (2016). Deep Reinforcement Learning with Double Q-learning. AAAI 2016, 2094-2100  
<http://arxiv.org/abs/1509.06461>
- Schaul, T., Quan, J., Antonoglou, I., Silver, D. (2016). Prioritized Experience Replay. ICLR 2016  
<https://arxiv.org/abs/1511.05952>
- Arulkumaran, K., Deisenroth, M. P., Brundage, M., & Bharath, A. A. (2017). A Brief Survey of Deep Reinforcement Learning. IEEE SPM Special Issue on Deep Learning for Visual Understanding  
<http://arxiv.org/abs/1708.05866>